

Bio 724: Data wrangling

Paul Magwene and Greg Wray

Real-world data is often messy

Data files you generate or will be given may...

- ▶ Be poorly organized
- ▶ Have missing values
- ▶ Contain extraneous information
- ▶ Lack headers (variable names)
- ▶ Confound variables and labels
- ▶ Use different encoding schemes
- ▶ Use unfamiliar conventions for dates, decimal separators, etc.
- ▶ Include empty columns – used for visual organization in spreadsheet, but interferes with analysis
- ▶ Include meta data and comments

Tidy data

To facilitate downstream analyses, data should be organized in a manner such that...

1. Each variable must have its own column.
2. Each observation must have its own row.
3. Each value must have its own cell.

country	year	cases	population
Afghanistan	1999	1845	15467071
Afghanistan	2000	1866	20535360
Brazil	1999	31737	172406362
Brazil	2000	81488	174404898
China	1999	211258	1272015272
China	2000	213766	128043583

variables

country	year	cases	population
Afghanistan	1999	1845	15467071
Afghanistan	2000	1866	20535360
Brazil	1999	31737	172406362
Brazil	2000	81488	174404898
China	1999	211258	1272015272
China	2000	213766	128043583

observations

country	year	cases	population
Afghanistan	1999	1845	15467071
Afghanistan	2000	1866	20535360
Brazil	1999	31737	172406362
Brazil	2000	81488	174404898
China	1999	211258	1272015272
China	2000	213766	128043583

values

Figure 1: Visual representation of tidy data (from R4DS2e).

dplyr and tidyr to the rescue

The tidyverse packages `dplyr` and `tidyr` provide many useful tools for wrangling data into a tidy form.

dplyr functions that facilitate wrangling

- ▶ `select`
- ▶ `filter`
- ▶ `mutate`
- ▶ `rename`
- ▶ `arrange`
- ▶ `left_join`
- ▶ `full_join`
- ▶ `inner_join`

key functions introduced by tidyr

- ▶ `pivot_longer` - reshape column data into rows
- ▶ `pivot_wider` - reshape row data into columns
- ▶ `separate_wider_delim`, `separate_wider_position`, `separate_wider_regex` - turns a single column into multiple columns
- ▶ `unite` - turns multiple columns in a single column

Example: Starting messy data

```
# Generated by Alexa and Brandon
,DrugX_low,DrugX_high,,DrugY_low_rep1,DrugY_low_rep2,DrugY_high_rep1,DrugY_high_rep2
GeneA,0.15,,, -0.15,-0.15,-0.21,-0.22
GeneB,-0.07,-0.76,, -0.11,0.1,0.01,-0.12
GeneC,-1.22,-0.27,, -0.14,,0.1,-0.1
GeneD,-0.09,1.2,, -0.02,-0.48,-0.11,0.16
GeneE,-0.6,1.01,, -0.05,-0.53,-0.47,0.24
GeneF,,1.39,,, -0.13,
```

	A	B	C	D	E	F	G	H
1	# Generated by Alexa and Brandon							
2		DrugX_low	DrugX_high		DrugY_low_rep1	DrugY_low_rep2	DrugY_high_rep1	DrugY_high_rep2
3	GeneA	0.15			-0.15	-0.15	-0.21	-0.22
4	GeneB	-0.07	-0.76		-0.11	0.1	0.01	-0.12
5	GeneC	-1.22	-0.27		-0.14		0.1	-0.1
6	GeneD	-0.09	1.2		-0.02	-0.48	-0.11	0.16
7	GeneE	-0.6	1.01		-0.05	-0.53	-0.47	0.24
8	GeneF		1.39				-0.13	
9								

Figure 2: In what ways is this data non-tidy?

Naively reading the data produces poor results

```
messy <- read_csv("~/Downloads/small-messy-data.csv")
```

```
# Generated by Alexa and Brandon
```

```
,DrugX_low,DrugX_high,,DrugY_low_rep1,DrugY_low_rep2,DrugY_high_rep1,DrugY_high_rep2
```

```
GeneA,0.15,, -0.15,-0.15,-0.21,-0.22
```

```
GeneB,-0.07,-0.76,, -0.11,0.1,0.01,-0.12
```

```
GeneC,-1.22,-0.27,, -0.14,,0.1,-0.1
```

```
GeneD,-0.09,1.2,, -0.02,-0.48,-0.11,0.16
```

```
GeneE,-0.6,1.01,, -0.05,-0.53,-0.47,0.24
```

```
GeneF
```

```
,1.39
```

```
,
```

```
,-0.13
```

Explore options of your reader function(s)

Example: filtering comment lines

```
messy <- read_csv("~/Downloads/small-messy-data.csv", comment="#")
```

...1	DrugX_low	DrugX_high	...4	DrugY_low_rep1	DrugY_low_rep2	DrugY_high_rep1	DrugY_high_rep2
GeneA	0.15	NA	NA	-0.15	-0.15	-0.21	-0.22
GeneB	-0.07	-0.76	NA	-0.11	0.10	0.01	-0.12
GeneC	-1.22	-0.27	NA	-0.14	NA	0.10	-0.10
GeneD	-0.09	1.20	NA	-0.02	-0.48	-0.11	0.16
GeneE	-0.60	1.01	NA	-0.05	-0.53	-0.47	0.24
GeneF	NA	1.39	NA	NA	NA	-0.13	NA

Renaming columns using dplyr::rename

```
messy <-  
  messy |>  
  rename(Gene = "...1")
```

Gene	DrugX_low	DrugX_high	...4	DrugY_low_rep1	DrugY_low_rep2	DrugY_high_rep1	DrugY_high_rep2
GeneA	0.15	NA	NA	-0.15	-0.15	-0.21	-0.22
GeneB	-0.07	-0.76	NA	-0.11	0.10	0.01	-0.12
GeneC	-1.22	-0.27	NA	-0.14	NA	0.10	-0.10
GeneD	-0.09	1.20	NA	-0.02	-0.48	-0.11	0.16
GeneE	-0.60	1.01	NA	-0.05	-0.53	-0.47	0.24
GeneF	NA	1.39	NA	NA	NA	-0.13	NA

Dropping columns using select

```
messy |>  
  select(-"...4")
```

Gene	DrugX_low	DrugX_high	DrugY_low_rep1	DrugY_low_rep2	DrugY_high_rep1	DrugY_high_rep2
GeneA	0.15	NA	-0.15	-0.15	-0.21	-0.22
GeneB	-0.07	-0.76	-0.11	0.10	0.01	-0.12
GeneC	-1.22	-0.27	-0.14	NA	0.10	-0.10
GeneD	-0.09	1.20	-0.02	-0.48	-0.11	0.16
GeneE	-0.60	1.01	-0.05	-0.53	-0.47	0.24
GeneF	NA	1.39	NA	NA	-0.13	NA

Dropping columns using select and where

If you had many columns it might not be feasible to specify the names directly. The `where` helper function can be used to specify a function to determine column selection.

```
messy <-  
  messy |>  
  select(-where( function(x) all(is.na(x)) ))
```

Gene	DrugX_low	DrugX_high	DrugY_low_rep1	DrugY_low_rep2	DrugY_high_rep1	DrugY_high_rep2
GeneA	0.15	NA	-0.15	-0.15	-0.21	-0.22
GeneB	-0.07	-0.76	-0.11	0.10	0.01	-0.12
GeneC	-1.22	-0.27	-0.14	NA	0.10	-0.10
GeneD	-0.09	1.20	-0.02	-0.48	-0.11	0.16
GeneE	-0.60	1.01	-0.05	-0.53	-0.47	0.24
GeneF	NA	1.39	NA	NA	-0.13	NA

Reshaping a data frame using pivoting

`tidyr::pivot_longer` collapses multiple columns into a single column, and create a new column from the collapse columns headers:

```
long_messy <-  
  messy |>  
  pivot_longer(cols = -Gene,  
               names_to="Drug_Dosage_Rep",  
               values_to="Expression")
```

Gene	Drug_Dosage_Rep	Expression
GeneA	DrugX_low	0.15
GeneA	DrugX_high	NA
GeneA	DrugY_low_rep1	-0.15
GeneA	DrugY_low_rep2	-0.15
GeneA	DrugY_high_rep1	-0.21
GeneA	DrugY_high_rep2	-0.22
GeneB	DrugX_low	-0.07
GeneB	DrugX_high	-0.76
GeneB	DrugY_low_rep1	-0.11
GeneB	DrugY_low_rep2	0.10

Extract a column in multiple columns

`tidyr::separate_wider_delim` splits a single column into multiple columns based on a character delimiter.

```
split_long_messy <-  
  long_messy |>  
  separate_wider_delim(cols = Drug_Dosage_Rep,  
                        delim = "_",  
                        names = c("Drug", "Dosage", "Replicate"),  
                        too_few = "align_start")
```

Gene	Drug	Dosage	Replicate	Expression
GeneA	DrugX	low	NA	0.15
GeneA	DrugX	high	NA	NA
GeneA	DrugY	low	rep1	-0.15
GeneA	DrugY	low	rep2	-0.15
GeneA	DrugY	high	rep1	-0.21
GeneA	DrugY	high	rep2	-0.22
GeneB	DrugX	low	NA	-0.07
GeneB	DrugX	high	NA	-0.76
GeneB	DrugY	low	rep1	-0.11
GeneB	DrugY	low	rep2	0.10

Filling missing data

`tidyr::replace_na` can be used to replace NA values with a default:

```
tidy_data <-  
  split_long_messy |>  
  replace_na(list(Replicate = "rep1"))
```

Gene	Drug	Dosage	Replicate	Expression
GeneA	DrugX	low	rep1	0.15
GeneA	DrugX	high	rep1	NA
GeneA	DrugY	low	rep1	-0.15
GeneA	DrugY	low	rep2	-0.15
GeneA	DrugY	high	rep1	-0.21
GeneA	DrugY	high	rep2	-0.22
GeneB	DrugX	low	rep1	-0.07
GeneB	DrugX	high	rep1	-0.76
GeneB	DrugY	low	rep1	-0.11
GeneB	DrugY	low	rep2	0.10


Tidy data enables complex summaries

```
tidy_summary <-  
  tidy_data |>  
  group_by(Gene, Drug, Dosage) |>  
  summarize(Mean_Expression = mean(Expression, na.rm=TRUE))
```

Gene	Drug	Dosage	Mean_Expression
GeneA	DrugX	high	NaN
GeneA	DrugX	low	0.150
GeneA	DrugY	high	-0.215
GeneA	DrugY	low	-0.150
GeneB	DrugX	high	-0.760
GeneB	DrugX	low	-0.070
GeneB	DrugY	high	-0.055
GeneB	DrugY	low	-0.005
GeneC	DrugX	high	-0.270
GeneC	DrugX	low	-1.220


Tidy data enables complex plotting

```
tidy_summary |>  
  ggplot(aes(x = Dosage, y = Mean_Expression, color=Drug)) +  
  geom_point(alpha=0.85, size = 3) +  
  geom_line(aes(group=Drug)) +  
  facet_wrap(~Gene)
```



Summary and raw values as different ggplot layers

```
tidy_summary |>
  ggplot(aes(x = Dosage, y = Mean_Expression, color=Drug)) +
  geom_point(alpha=0.85, size=3) +
  geom_line(aes(group=Drug)) +
  geom_point(data=tidy_data,
             mapping=aes(x = Dosage, y = Expression, color=Drug),
             alpha = 0.5, size = 1.5,
             inherit.aes = FALSE) +
  facet_wrap(~Gene)
```



Widening

`tidyr::pivot_wider` creates new columns and headers from specified columns:

```
tidy_wide <-  
  tidy_summary |>  
  pivot_wider(names_from = Gene, values_from = Mean_Expression)
```

Drug	Dosage	GeneA	GeneB	GeneC	GeneD	GeneE	GeneF
DrugX	high	NaN	-0.760	-0.27	1.200	1.010	1.39
DrugX	low	0.150	-0.070	-1.22	-0.090	-0.600	NaN
DrugY	high	-0.215	-0.055	0.00	0.025	-0.115	-0.13
DrugY	low	-0.150	-0.005	-0.14	-0.250	-0.290	NaN

```
corr_from_wide <-  
  tidy_summary |>  
  ungroup() |> # ungroup important here  
  pivot_wider(names_from = Gene, values_from = Mean_Expression) |>  
  select(-c(Drug,Dosage,GeneF)) |>  
  cor(use = "pairwise.complete.obs" )
```

	GeneA	GeneB	GeneC	GeneD	GeneE
GeneA	1.0000000	-0.5464173	-0.9980371	-0.0734459	-0.9807466
GeneB	-0.5464173	1.0000000	-0.1197536	-0.9929120	-0.9472035
GeneC	-0.9980371	-0.1197536	1.0000000	0.1652294	0.4313106
GeneD	-0.0734459	-0.9929120	0.1652294	1.0000000	0.9572841
GeneE	-0.9807466	-0.9472035	0.4313106	0.9572841	1.0000000